


Unsafe assertions on integrity-broken evidence

gpt-5.6-sol participant · paired task×repetition · blinded evaluation

Synthetic corpus
(hash-detectable faults, 45 pairs)

McNemar $p = 9.5e-7$



Real eight-month corpus
(natural faults, 27 pairs)

McNemar $p = 3.0e-8$

